

notebook

September 8, 2025

1 Import Libraries

```
[1]: import functools
import tracemalloc
import psutil
import os
from pathlib import Path
import time

import httpx
import pandas as pd
import seaborn as sns
from matplotlib import pyplot as plt
```

2 Utilities

```
[2]: def profiler(func):
    """Decorator to measure memory usage and execution time of a function."""
    @functools.wraps(func)
    def wrapper(*args, **kwargs):
        process = psutil.Process(os.getpid())

        # Start memory + time tracking
        start_mem = process.memory_info().rss / 1024**2
        tracemalloc.start()
        start_time = time.time()

        result = func(*args, **kwargs) # run target function

        # After execution
        current, peak = tracemalloc.get_traced_memory()
        end_mem = process.memory_info().rss / 1024**2
        end_time = time.time()
        tracemalloc.stop()

        print(f"\n--- Memory Profile for `{func.__name__}` ---")
        print(f"Start memory    : {start_mem:.2f} MB")
```

```

print(f"End memory      : {end_mem:.2f} MB")
print(f"Peak (tracked)  : {peak / 1024**2:.2f} MB")
print(f"Execution time  : {end_time - start_time:.2f} sec")
print("-----\n")

return result
return wrapper

```

```

[3]: # Constants

ROOT_PATH = Path(os.getcwd())
DATASET_URL = "https://drive.usercontent.google.com/download?
↳id=1N1xoxgcw2K3d-49tlchXAWw4wuxLj7EV&export=download"
DATASET_OUTPUT_PATH = ROOT_PATH / "dataset.csv"

```

```

[4]: # Utilities
@profiler
def download_data(url: str, output_path: Path) -> None:
    with httpx.stream("GET", url) as response:
        response.raise_for_status() # check for HTTP errors
        with open(output_path, "wb") as f:
            for chunk in response.iter_bytes():
                f.write(chunk)

    print(f"Downloaded to {output_path}")

@profiler
def load_data(file_path: Path, **kwargs) -> pd.DataFrame:
    return pd.read_csv(file_path, **kwargs)

```

3 EDA

3.1 Data Loading

```

[ ]: # Data Loading
download_data(DATASET_URL, DATASET_OUTPUT_PATH)

```

```

[13]: # Data Parsing
df = load_data(DATASET_OUTPUT_PATH, index_col=0)

```

```

--- Memory Profile for `load_data` ---
Start memory   : 193.36 MB
End memory     : 304.31 MB
Peak (tracked) : 56.09 MB
Execution time  : 1.75 sec
-----

```

3.2 Data Understanding

```
[6]: df.shape
```

```
[6]: (100000, 11)
```

```
[7]: df.head()
```

```
[7]:
```

	Customer Id	First Name	Last Name	Company \
Index				
1	ffeCAb7AbcB0f07	Jared	Jarvis	Sanchez-Fletcher
2	b687FfC4F1600eC	Marie	Malone	Mckay PLC
3	9FF9ACbc69dcF9c	Elijah	Barrera	Marks and Sons
4	b49edDB1295FF6E	Sheryl	Montgomery	Kirby, Vaughn and Sanders
5	3dcCbFEB17CCf2E	Jeremy	Houston	Lester-Manning

	City	Country \
Index		
1	Hatfieldshire	Eritrea
2	Robertsonburgh	Botswana
3	Kimbury	Barbados
4	Briannaview	Antarctica (the territory South of 60 deg S)
5	South Brianna	Micronesia

	Phone 1	Phone 2 \
Index		
1	274.188.8773x41185	001-215-760-4642x969
2	283-236-9529	(189)129-8356x63741
3	8252703789	459-916-7241x0909
4	425.475.3586	(392)819-9063
5	+1-223-666-5313x4530	252-488-3850x692

	Email	Subscription Date \
Index		
1	gabriellehartman@benjamin.com	2021-11-11
2	kstafford@sexton.com	2021-05-14
3	jeanettecross@brown.com	2021-03-17
4	thomassierra@barrett.com	2020-09-23
5	rubenwatkins@jacobs-wallace.info	2020-09-18

	Website
Index	
1	https://www.mccarthy.info/
2	http://www.reynolds.com/
3	https://neal.com/

```
4      https://www.powell-bryan.com/
5      https://www.carrillo.com/
```

```
[8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 100000 entries, 1 to 100000
Data columns (total 11 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Customer Id           100000 non-null object
 1   First Name            100000 non-null object
 2   Last Name             100000 non-null object
 3   Company               100000 non-null object
 4   City                  100000 non-null object
 5   Country               100000 non-null object
 6   Phone 1               100000 non-null object
 7   Phone 2               100000 non-null object
 8   Email                 100000 non-null object
 9   Subscription Date     100000 non-null object
10   Website               100000 non-null object
dtypes: object(11)
memory usage: 9.2+ MB
```

```
[16]: # Checking missing values
df.isnull().sum()
```

```
[16]: Customer Id           0
      First Name          0
      Last Name           0
      Company             0
      City                0
      Country             0
      Phone 1             0
      Phone 2             0
      Email               0
      Subscription Date    0
      Website             0
      dtype: int64
```

3.3 Data Cleaning and Preparation

```
[15]: # Parsing to date
@profiler
def parse_date(df: pd.DataFrame, col: str):
    df[col] = pd.to_datetime(df[col])

parse_date(df, "Subscription Date")
```

```
--- Memory Profile for `parse_date` ---
Start memory   : 242.02 MB
End memory     : 242.78 MB
Peak (tracked) : 1.53 MB
Execution time : 0.10 sec
-----
```

```
[11]: df.nunique()
```

```
[11]: Customer Id      100000
      First Name      690
      Last Name       1000
      Company         71994
      City            49154
      Country         243
      Phone 1         100000
      Phone 2         100000
      Email           99995
      Subscription Date      880
      Website         50471
      dtype: int64
```

3.4 Descriptive Analysis

```
[130]: @profiler
def plot_top_n(src, col: str = None, title: str = None, top: int = 10):
    if isinstance(src, pd.DataFrame):
        if col is None:
            raise ValueError("When src is a DataFrame, you must provide a column name.")
        series = src[col]
    elif isinstance(src, pd.Series):
        series = src
    else:
        raise TypeError("src must be a pandas DataFrame or Series")

    ax = sns.countplot(
        y=series,
        order=series.value_counts().head(top).index
    )
    ax.set_title(title)
```

3.4.1 Company

```
[131]: plot_top_n(df, 'Company', "Top 10 Company Subscriber")
```

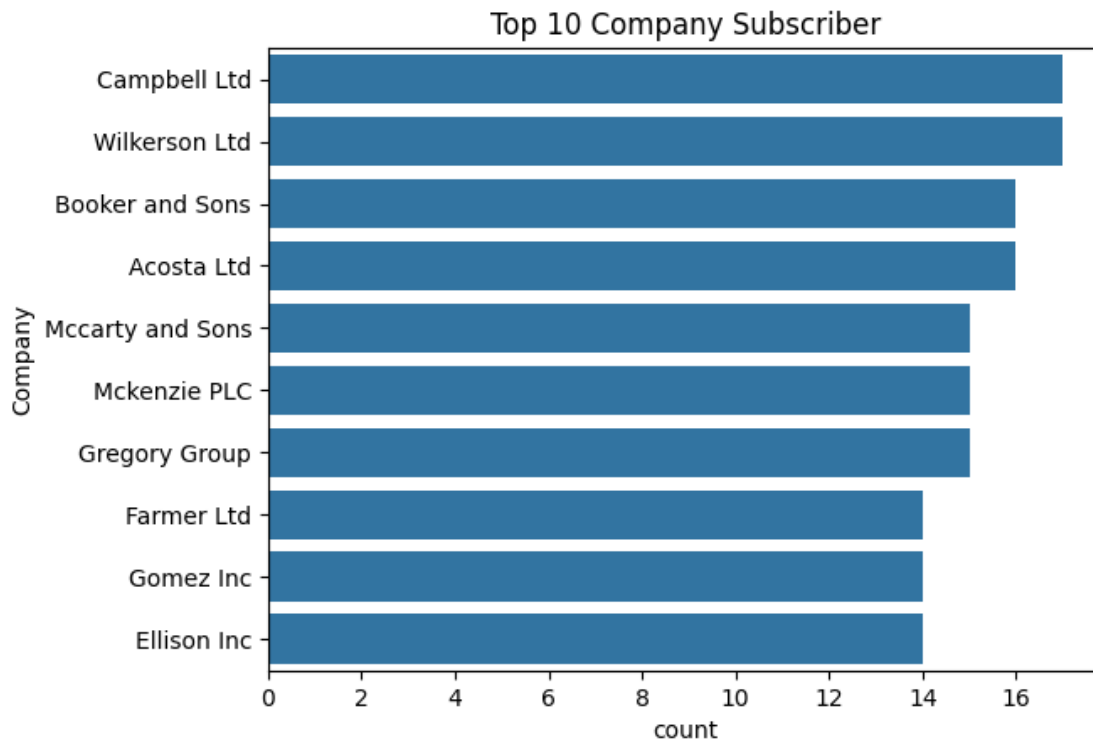
--- Memory Profile for `plot_top_n` ---

Start memory : 377.68 MB

End memory : 385.97 MB

Peak (tracked) : 11.51 MB

Execution time : 0.73 sec



```
[93]: @profiler
def plot_top_company_by_type(src):
    company_suffix = src["Company"].str.split().apply(lambda x: x[-1])
    plot_top_n(company_suffix, title="Top 10 Company Subscribers By Company_
↳Type")

plot_top_company_by_type(df)
```

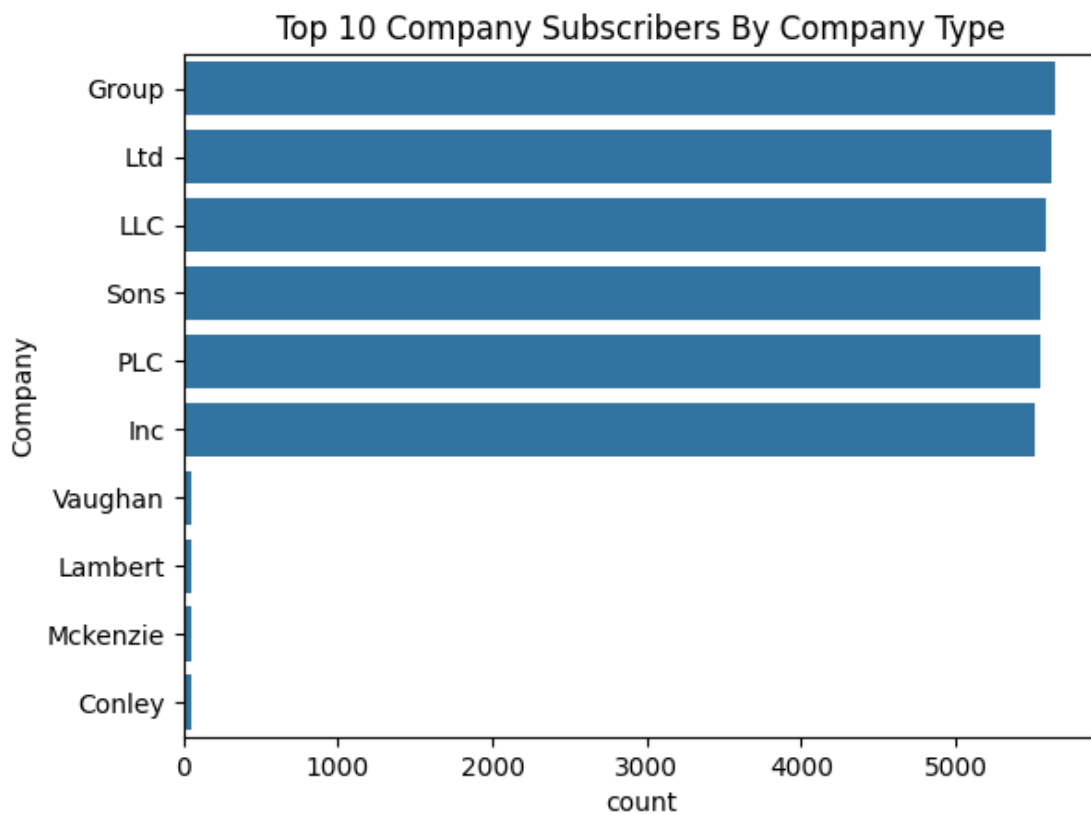
--- Memory Profile for `plot_top_10` ---

Start memory : 384.72 MB

```
End memory      : 392.74 MB
Peak (tracked)  : 30.62 MB
Execution time  : 1.07 sec
-----
```

--- Memory Profile for `plot_top_company_by_type` ---

```
Start memory    : 377.88 MB
End memory      : 358.63 MB
Peak (tracked)  : 0.00 MB
Execution time  : 3.04 sec
-----
```



3.4.2 Subscription Date

```
[94]: @profiler
def plot_counts_over_time(df: pd.DataFrame, date_col: str, count_col: str):
    # Ensure datetime
    df[date_col] = pd.to_datetime(df[date_col])
```


```

# Aggregate counts per date
df_counts = df.groupby(date_col)[count_col].count().reset_index()

# Plot
sns.lineplot(data=df_counts, x=date_col, y=count_col)
plt.xticks(rotation=45)
plt.xlabel("Date")
plt.ylabel("Count")
plt.title(f"Counts of {count_col} over time")
plt.show()

plot_counts_over_time(df, "Subscription Date", "Customer Id")

```



```

--- Memory Profile for `plot_counts_over_time` ---
Start memory   : 360.22 MB
End memory     : 365.00 MB
Peak (tracked) : 3.57 MB

```

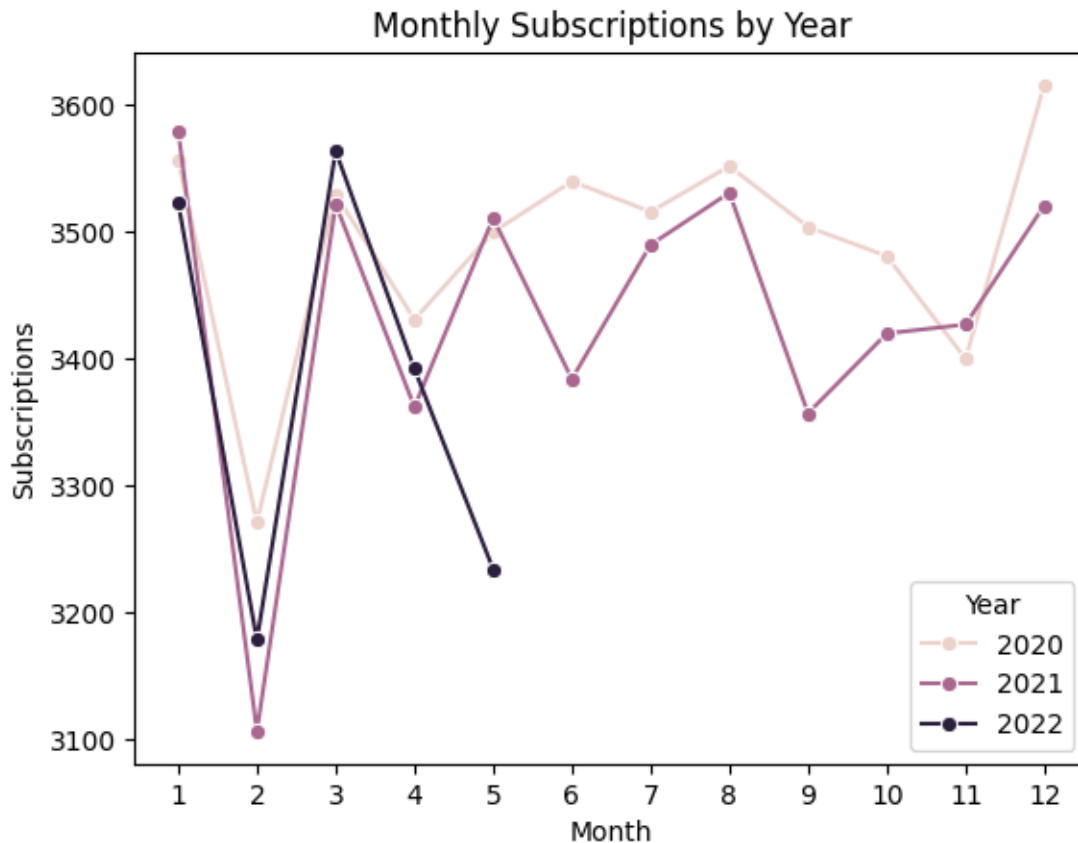

Execution time : 2.38 sec

```
[147]: @profiler
def plot_subscriptions_monthly(df: pd.DataFrame, date_col: str):
    df["Year"] = df[date_col].dt.year
    df["Month"] = df[date_col].dt.month

    monthly_counts = df.groupby(["Year", "Month"]).size().
↪reset_index(name="Count")
    sns.lineplot(
        data=monthly_counts,
        x="Month",
        y="Count",
        hue="Year",      # separate line per year
        marker="o"       # dots on points (optional)
    )

    plt.xticks(range(1, 13)) # months 1-12
    plt.xlabel("Month")
    plt.ylabel("Subscriptions")
    plt.title("Monthly Subscriptions by Year")
    plt.show()

plot_subscriptions_monthly(df, "Subscription Date")
```




```
--- Memory Profile for `plot_subscriptions_monthly` ---
Start memory   : 424.00 MB
End memory     : 426.63 MB
Peak (tracked) : 6.70 MB
Execution time : 1.54 sec
-----
```

3.4.3 Country

```
[97]: plot_top_10(df, "Country", title="Top 10 Countries By Subscribers")
```

```
--- Memory Profile for `plot_top_10` ---
Start memory   : 368.26 MB
End memory     : 371.79 MB
Peak (tracked) : 10.96 MB
Execution time : 0.74 sec
-----
```

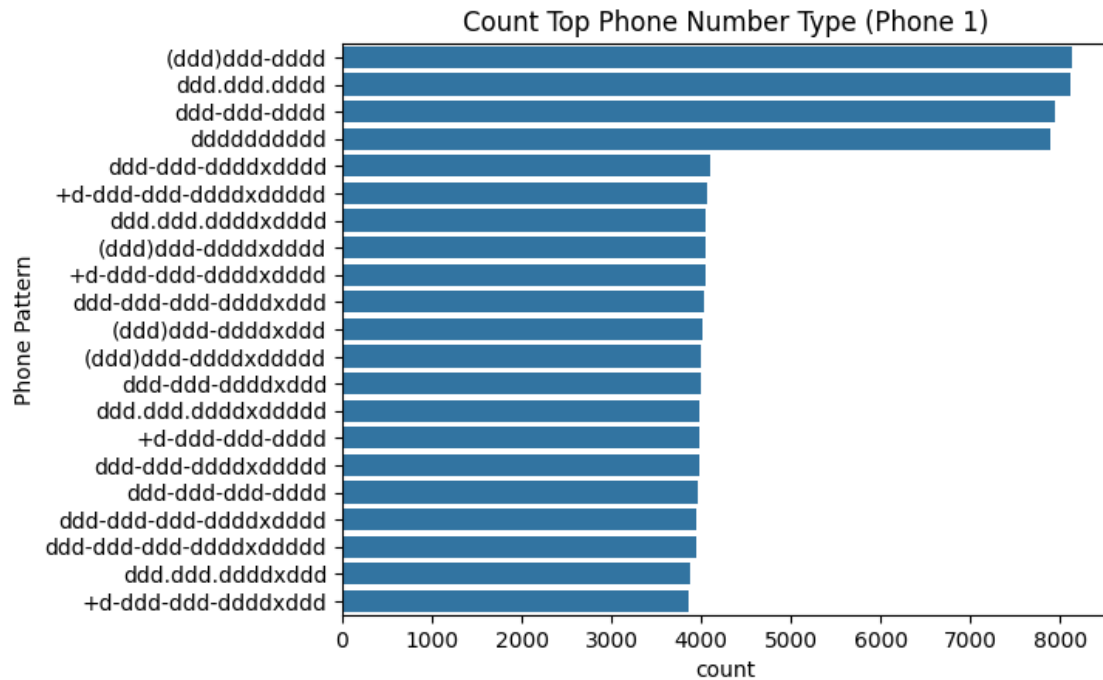


3.4.4 Phone Number

```
[141]: def plot_phone_number_pattern(df: pd.DataFrame, col: str):
        df['Phone Pattern'] = df[col].str.replace(r"\d", "d", regex=True)
        n = df['Phone Pattern'].nunique()
        plot_top_n(df, 'Phone Pattern', title=f"Count Top Phone Number Type_
↪({col})", top=n)
```

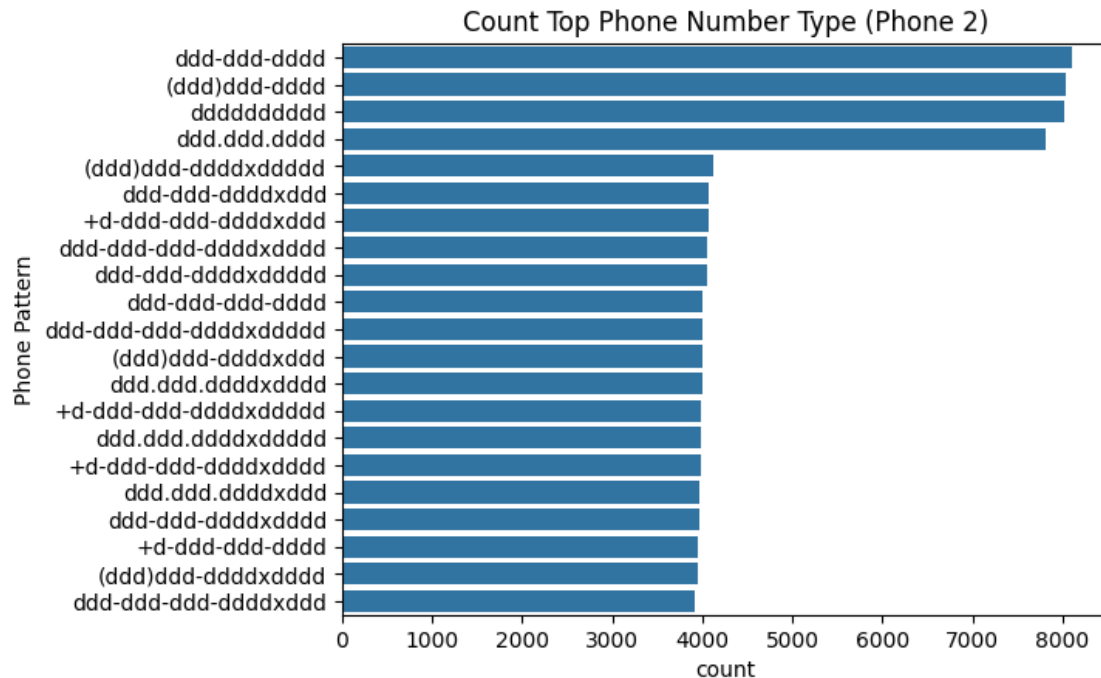
```
[143]: plot_phone_number_pattern(df, "Phone 1")
```

```
--- Memory Profile for `plot_top_n` ---
Start memory   : 405.56 MB
End memory     : 409.86 MB
Peak (tracked) : 11.81 MB
Execution time  : 1.09 sec
-----
```



```
[144]: plot_phone_number_pattern(df, "Phone 2")
```

```
--- Memory Profile for `plot_top_n` ---
Start memory   : 416.24 MB
End memory     : 423.22 MB
Peak (tracked) : 11.81 MB
Execution time : 1.07 sec
```



3.4.5 Email

```
[157]: email_domain = df["Email"].str.split("@").apply(lambda x: x[-1])
       email_domain.nunique()
```

```
[157]: 38322
```

```
[154]: plot_top_n(email_domain, title="Top 10 Company Subscribers By Company Type")
```

```
--- Memory Profile for `plot_top_n` ---
```

```
Start memory   : 385.48 MB
```

```
End memory     : 395.90 MB
```

```
Peak (tracked) : 11.25 MB
```

```
Execution time : 0.77 sec
```

```
-----
```

